



Exercises for **Programming, Data Analysis, and Deep Learning in Python** (SoSe 2021)

Exercise Sheet no. 11, *Deadline*: Monday, July 5, 10:15

Notes

- Pay attention to the notes on the previous sheet.

Exercise 31 Polynomial Regression with Least Squares (programming ex.) (12 points)

Download the corresponding Python script from E-Learning. The script contains data points (x_i, y_i) , $i = 1, \dots, m$. The goal is to approximate the given data set by a polynomial of degree $n < m$ using the Least Squares method. To this end, complete the following tasks:

- 1) For a given degree n , create either of the following two matrices (m rows, $n+1$ columns or $n+1$ rows, m columns) using the x -coordinate of the respective data point:

$$A := \begin{pmatrix} 1 & x_1 & x_1^2 & \dots & x_1^n \\ \vdots & \vdots & \vdots & \vdots & \vdots \\ 1 & x_m & x_m^2 & \dots & x_m^n \end{pmatrix} \quad \text{or} \quad A^\top := \begin{pmatrix} 1 & \dots & 1 \\ x_1 & \dots & x_m \\ x_1^2 & \dots & x_m^2 \\ \vdots & \dots & \vdots \\ x_1^n & \dots & x_m^n \end{pmatrix}.$$

- 2) Calculate the matrix $A^\top A$ and store the result in the variable `ATA`. From the y -coordinates of the data points, construct the vector $b := A^\top y$.
- 3) Solve the system of linear equations $A^\top A a = b$ for a in two different ways:
 - directly, using `np.linalg.solve`, and
 - by calculating the inverse of $A^\top A$, and store the results in variables called `a_direct` and `a_via_inv`, respectively. (The vector $a = (a_0, a_1, \dots, a_n)$ contains the coefficients of the approximating polynomial.)
- 4) Plot the data points (x_i, y_i) as single (unconnected) red dots and also display both approximating polynomials in that plot.
Hint: To plot a polynomial $p(x)$, first, create a grid `x_fine` with `np.arange(...)` for suitable parameters. Then evaluate the polynomial on that grid (you may use the function `f_eval` for that) and store the result in a list or array `f_x`. Finally, use `x_fine` and `f_x` in the `plt.plot()` command to plot the polynomial.
- 5) Calculate the errors $\|Aa - y\|_2$ for $a = \text{a_direct}$ and $a = \text{a_via_inv}$.

Try out different polynomial degrees $n = 1, \dots, 15$. Does the error for `a_direct` always decrease with higher n ? Briefly describe your observations regarding the error for `a_via_inv`.

Exercise 32 Predicting Diabetes (programming exercise)

(12 points)

Download the associated file¹ from E-Learning. In this exercise we will try to predict whether a person has Diabetes using logistic regression.

- a) Pre-process the data as follows. First, read the csv file, then divide the columns into two types of variables: The *target variable* (also called *dependent variable*) is the last column, which you should store in the variable `y`. The *feature variables* (also called *independent variables*) are all other columns. Store these in the variable `X`.
Next, as in the lecture, divide `X` and `y` into training data (`train_set_x`, `train_set_y`) and test data (`test_set_x`, `test_set_y`) using 75% of the data as training data and the other 25% as test data.
- b) Copy all the functions that you need for logistic regression from the lecture notes (including `initialize_parameters(dim)`) and modify the following: In the `propagate` function, after the `cost` is calculated, check whether it is `NaN` and if so, change it to `np.inf`. In the `optimize` function, every 10000 steps (instead of every 100 steps), append the `cost` to the `costs` list and output the quadratic Euclidean norm of the gradient of the cost: `np.sum(grads['dw'] ** 2) + np.sum(grads['db'] ** 2)`
- c) As in the lecture notes, call the `model` function to run the logistic regression and then plot the costs.
Important: Set `np.random.seed(0)` before the pre-processing step. When calling the `model` function, set the number of steps to 1000001 and the learning rate to 0.00025.
Hints: When calling the `model` function, you might need to fix any errors regarding the dimension of arrays. You can do so by comparing the shapes of the respective arrays to the ones from the lecture.
- d) Output the so-called *confusion matrix*. It displays the number of correct predictions on the diagonal and the incorrect predictions on the off-diagonal, similar to:

		Predicted	
		0	1
Actual	0	118	12
	1	26	36

To output it, proceed as follows: First, import `metrics` from the `sklearn` module: `from sklearn import metrics`. Then use `predict` to calculate the predictions on the test set. Next, call `cnf_matrix = metrics.confusion_matrix(arg1, arg2)`, where you should replace `arg1` by the test set of `y` and `arg2` by the predictions on that test set. Finally, `print(cnf_matrix)`.

- e) Perform the tasks c) and d) with
 - a learning rate of 0.00025, once with 100001 steps and once with 10001 steps,
 - a learning rate of 0.0002, once with 1000001 steps, once with 100001 steps, and once with 10001 steps,
 and compare the corresponding confusion matrices.

¹Source (2020-01-14): <https://github.com/plotly/datasets/blob/master/diabetes.csv>